

Optimizing For Low Latency

Ramakrishna Karedla
Financial Services Lab
Parsippany, NJ



The Challenges to Nano-Second Latencies

Network

Data transfer on the network,
Network Hardware Latencies, Network Drivers Latencies, etc

Software

Sync overhead and oversubscription of threads,
Interrupts, I/O, Parallel Programming, Software Tools, etc

Compute

Frequency, Instructions per cycle,
Cache Sizes, Memory Speed & Bandwidth, RAS



Achieving Low Latency

Maximizing Performance involves “Hardware & System Level” optimizations

- Implement low latency BIOS settings recommended by the OEM
 - OEM BIOS Settings: SMIs, HyperThreading, C-States- All Off.
 - Better Turbo stability in SandyBridge over previous version. Recommend enabling
 - Turn Memory Power Savings Off.
- On the application side: Maximize your resources by...
 - Pin Threads, Interrupts, and Processes to individual cores . CPUsets, isolcpus
 - Use the DPDK tool kit to optimize code www.intel.com/go/dpdk
 - Place “communicating” threads on adjacent cores. Avoid Global variables and partial cache line writes
 - Avoid use of locking algorithms; Keep retry section to at least 500 nsecs
<http://www.intel.com/content/dam/www/public/us/en/documents/white-papers/xeon-lock-scaling-analysis-paper.pdf>
 - Avoid mixing SSE instructions with AVX instructions in same code segment
use separate compilation units or use vzeroupper
 - Evaluate IPP / SSE 4.2 for string comparison functions
 - ✓ <http://software.intel.com/en-us/articles/using-avx-without-writing-avx-code>



Achieving Low Latency (contd)

- Determine how many cores your trading strategy requires. Avoid cross socket penalties. Match CPU+NIC as per DDIO strategy. Use polling to read data from L3 cache before data gets flushed !
- If using Red Hat, refer Red Hat suggestions for low latency profiles
<https://access.redhat.com/site/articles/221153>
- If using acpi-cpufreq driver, suggest using performance scaling governor
✓ <https://access.redhat.com/knowledge/articles/221153>
- MOVBE new instruction in Haswell. Suitable for network packet format conversions instead of BSWAP to reduce latency. Can use PSHUFB on SandyBridge
- Reduce kernel jitter by minimizing Ring Transitions. E.g Use NOPs (0.5 cycle) versus nanosleep()
- Read assembly (not difficult !). Compiler can always optimize code away !



Achieving Low Latency (contd)

- Linux, For C0 states Use idle=mwait instead of idle=poll. Same results. Less power.
Windows, use `powercfg.exe /setacvalueindex SCHEME_CURRENT SUB_PROCESSOR 5d76a2ca-e8c0-402f-a133-2158492d58ad 1`
`powercfg.exe /setactive SCHEME_CURRENT`
- Use RDTSCP instead of `gettimeofday()` to profile short code sections.
RDTSCP invariant across cores and sockets in a box (with no node controllers)
- For specific code sections, use `CPUID` to serialize instructions to avoid Branch mispredictions that may cause jitter.
- Optimize NIC driver settings for low latency rather throughput
- At times, SystemTap scripts can give a better insight into kernel behavior
<http://sourceware.org/systemtap/>
- VTune can now profile Java code. Consider using jhiccup to observe jitter
<http://jhiccup.com/>
- If business case permits, consider using Intel Xeon PHI for low latency workloads
<http://software.intel.com/mic-developer>

VTune™ Amplifier XE 2013 - Power Analysis Functionality

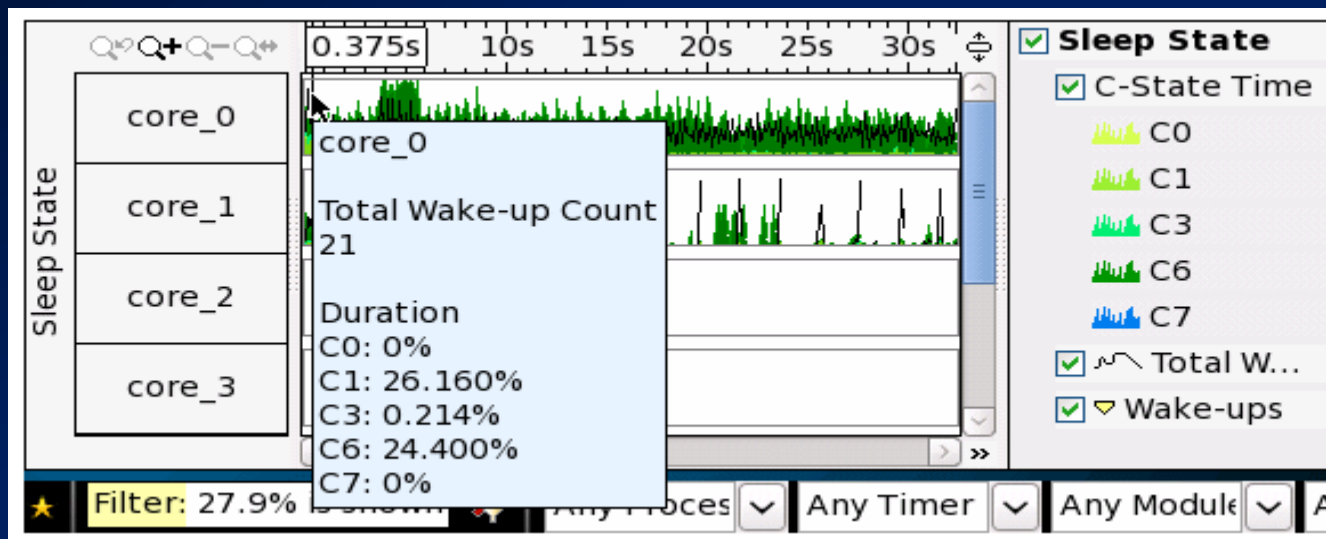
CPU Frequency & Sleep State Analysis - Timeline

CPU Frequency - CPU Frequency

Analysis Target Analysis Type Collection Log Summary

Grouping: Core

Core	Time				
	1.596GHz	1.729GHz	1.995GHz	2.661GHz	2.66GHz
core_0	0.541s	0s	0s	55.537s	0s
core_1	56.039s	0s	0s	0.040s	0s
core_5	56.078s	0s	0s	0.000s	0s
core_4	56.078s	0s	0s	0.000s	0s
core_3	56.075s	0s	0s	0.000s	0.003s
core_2	56.078s	0s	0s	0.000s	0s



Intel's Performance Counter Monitor – Real Time Tool

- Complements Intel's Flagship Profiler: Vtune Amplifier
- Can be used stand alone or embedded in code segments
- SOCH0; DRAMClocks: 665224950; Rank0 CKE Off Residency: 0.00%
- Core Frequency transition count: 0 (useful to study Outliers)

•	Core (SKT)	IPC	FREQ	AFREQ	L3MISS	L2MISS	L3HIT	L2HIT	L3CLK	L2CLK	TEMP	
•	0	0	0.25	1.00	1.15	802	1557	0.48	0.45	0.00	0.00	56
•	1	0	0.25	1.00	1.15	3554	4671	0.24	0.38	0.00	0.0	62
•	2	0	0.25	1.00	1.15	49	608	0.92	0.45	0.00	0.00	60
•	3	0	0.25	1.00	1.15	86	312	0.72	0.28	0.00	0.00	61
•	4	1	0.25	1.00	1.15	805	6764	0.88	0.44	0.00	0.00	64

C0 (active,non-halted) core residency: 100.00 %

C1 core residency: 0 %; C3 core residency: 0 %; C6 core residency: 0 %; C2 package residency: 0 %; C3 package residency: 0 %; C6 package residency: 0 %

Also shows QPI Traffic and Memory Read /Write Performance Monitoring
SMI count via Turbostat (MSR 0x34)



Financial Services Lab - NJ

Provides testing of new H/W and Software technologies to financial vertical

- Optimization of ISV and end user applications (banks, trading and market data firms on pre release & latest hardware: Processors, NICs, Switches etc & Intel's software tools
- Collaborate with OEMs and Industry benchmarks such as STAC T (Tick to Trade Latency), STAC N, STAC A2, STAC M3 etc
- Dissemination of non confidential best practices via <http://financialservices.intel.com>
- Investigations into Precision Time Protocol (PTP -1588) for clock synch in cluster.
- Recent STAC T (partner with STAC, Corvil, Cisco, Mellanox, Redline ISV, Dell, Datacom) highlighting Dell's DPAT Technology. A 13 % to 42 % reduction in latency over an earlier benchmark.

Market Data	Mean	Median	99 th %	Max	Std Dev
emini 1x	5.4	5.0	10.7	16.2	1.3
emini 8x	5.2	4.8	10.5	16.8	1.2

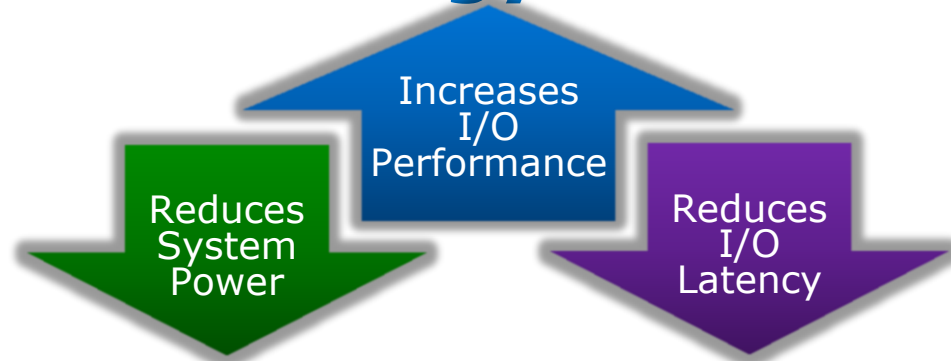
Backup



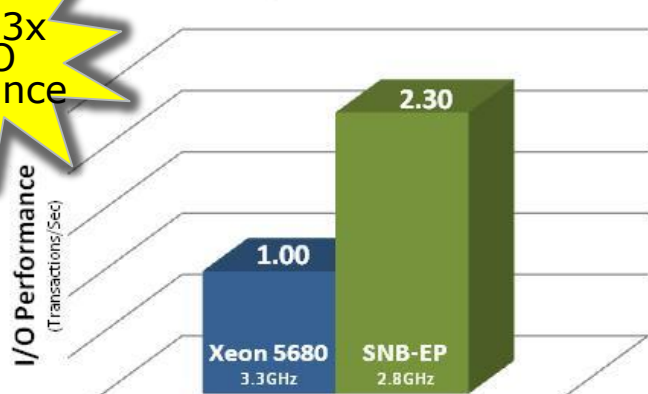
Intel® Data Direct I/O Technology

(Intel® DDIO)

- New Romley I/O architecture that leverages Intel® Integrated I/O
 - PCIe* lanes on the CPU
- Reduces memory accesses from I/O on local socket
 - Speeds up CPU data transfer
 - Accelerates inbound & outbound flows



SNB-EP I/O Performance




Up to 2.3x the I/O performance


Reduces I/O Accesses to Memory

DIMMs

Socket 0



Socket 1



QPI

QPI



Intel® Integrated I/O

- Romley Platform Ingredients
 - Sandy Bridge CPU (Skt-R or Skt-B2)
 - Intel Ethernet, PCIe* I/O devices

INTEL CONFIDENTIAL

New Instructions in Haswell


ISA Spec is public <http://software.intel.com/en-us/avx/>

Group	Description
Intel® AVX2	Adding vector integer operations to 256-bit
FMA	Fused Multiply-Add operation forms
Gather	Load elements using a vector of indices, vectorization enabler
RTM	Restricted Transactional Memory
Bit Manipulation and Cryptography	15 instructions improving performance of bit stream manipulation and decode, large integer arithmetic and hashes
MOVBE	Load and Store of Big Endian forms (previously introduced in Atom)
INVPCID	Invalidate processor context ID (Ring 0 instruction)

**This presentation covers vector ISA additions to Haswell;
other NIs covered in separate SES sessions**

Sandy Bridge “Ring Bus” Optimizes Platform Data Movement

- Ring bus architecture delivers an efficient bi-directional highway for data movement
- Compared to Xeon 5500/5600:
 - Lower L3 cache latency (~20%)
 - Higher bandwidth between cores, L3, Memory & I/O
 - Up to 8x more L3 to core bandwidth



Higher performance starts with The Ring!